

Lecture 18

March 29, 2022

Instructor: Sepehr Assadi

Disclaimer: *These notes have not been subjected to the usual scrutiny reserved for formal publications. They may be distributed outside this class only with the permission of the Instructor.*

1 Algorithms for MST

Recall that in the previous lecture, we provided the following generic “meta-algorithm” for finding MSTs:

1. Let $F = \emptyset$ be an empty *forest* initially. For $i = 1$ to $n - 1$ steps:
 - (a) Find a *safe* edge $e \in G - F$.
 - (b) Add e to F , namely, $F \leftarrow F \cup \{e\}$.
2. Return F as an MST of G .

Also recall the following two definitions:

- **MST-good forest:** We say that a forest F is *MST-good*, if there exists *some* MST T of G such that every edge of F belongs to T , namely, $F \subseteq T$ (note that MSTs are not unique in general).
- **Safe edge:** An edge e is called *safe* with respect to a forest F if the following condition holds: if F is MST-good, then $F \cup \{e\}$ is also MST-good.

We proved that as long as we can *implement* this meta-algorithm with an actual algorithm, we will get an algorithm for finding MSTs. The next step was to prove the following theorem:

Theorem 1. *Let $G = (V, E)$ be any undirected connected graph and F be any MST-good forest of G which is not a tree. Suppose $(S, V - S)$ is any cut in F with no cut edges. Then the edge e with minimum weight among the cut edges of $(S, V - S)$ in G is a safe edge with respect to F .*

We now only need to find a way of finding a minimum weight edge in some cut that is still “empty” in F . Both the algorithms we study in this course for MST, namely, Kruskal’s and Prim’s algorithms, exactly do this although via different approaches. In the rest of this lecture, we will see these two algorithms.

1.1 Kruskal’s Algorithm

Kruskal’s algorithm works as follows:

1. Sort the edges of G in increasing (non-decreasing) order of their weights.
2. Let $F = \emptyset$.
3. For $i = 1$ to m (in the sorted ordering of edges):
 - (a) If adding e_i to F does *not* create a cycle, let $F \leftarrow F \cup \{e_i\}$.
4. Return F .

We now prove the correctness of this algorithm and see how to implement the step of this algorithm that needs to check whether e_i creates a cycle or not in F efficiently, using the *union-find* data structure.

Proof of Correctness: Consider the forest F maintained by the algorithm. We prove that if F is MST-good at some iteration i and in that iteration we inserted an edge e_i to F , then e_i was safe with respect to F . This then implies that we only added safe edges to F . Moreover, the output of this algorithm is always a tree since whenever we see an edge that does not create a cycle we add it to F (and since G was connected, we will find a tree eventually this way). That means that we added $n - 1$ safe edges. This would imply the correctness of the algorithm as we now can say that this algorithm is an implementation of the meta-algorithm discussed in the previous lecture which we already proved its correctness.

Consider the iteration i in which we inserted the edge e_i to F . Let $e_i = \{u, v\}$. Since adding e_i to F did *not* create a cycle, we know that u and v were *not* in the same connected component of F . We can now define the cut $(S, V - S)$ where S is the connected component of u in F . Firstly, we know that none of the cut edges of G in $(S, V - S)$ belong to F (otherwise S would not have been a connected component of F). By discussion above we know that e_i is a cut edge of G in $(S, V - S)$ because $u \in S$ and $v \in V - S$. Finally, we know that e_i has the minimum weight among all cut edges of $(S, V - S)$: this is because at this point we know that e_i is the first cut edge of $(S, V - S)$ that we have visited in the algorithm (otherwise, we would have picked that edge sooner as it will not create a cycle in F , thus making it impossible for S to be a connected component of F); since the edges are sorted in non-decreasing order of their weights, we have that e_i has the minimum weight among cut edges of $(S, V - S)$.

We are now done since we can apply Theorem 1 and argue that e_i is a safe edge. This finalizes the proof.

Runtime Analysis: The first line of the algorithm takes $O(m \log m)$ time to sort all the edges (using, say, merge sort). We then have m iteration, each iteration involve deciding whether adding the edge e_i to the graph creates a cycle or not. The easiest way to implement this step is to run DFS/BFS – that will take $O(n + m)$ time per each iteration, making the total runtime of the algorithm $O(m \cdot (n + m)) = O(mn + m^2) = O(m^2)$ since we know that $m \geq n - 1$ as G was connected (any connected graph needs at least $n - 1$ edges).

However, this is *too inefficient* and we can in fact implement this line, using proper data structures called *union-find* (which we will describe below), in only $O(\log n)$ time with a preprocessing of $O(n)$ time – that will make the total runtime $O(n + m \log m + m \log n) = O(m \log m)$ (as again $m \geq n - 1$). It is worth pointing out that we can implement that step *much faster* almost (but not quietly) in constant time; however that will not reduce the runtime as we still need to sort the edges in $O(m \log m)$ time in general.

Thus, overall, the runtime of the algorithm is $O(m \log m)$.

Detour: Union-Find Data Structure

We now describe the union-find data structure, also called disjoint-set data structure, or merge-set data structure. The goal of this data-structure is to maintain a collection of *disjoint* subsets S_1, \dots, S_k from a universe $U := \{1, \dots, n\}$, allows us to find for each element $e \in U$, which of these sets e belong to, and to *merge*, or *union*, any two given sets S_i, S_j together, namely, remove S_i, S_j and replace them with $S_i \cup S_j$.

We now define the data structure formally. To do so, we simply need to define the operations supported by the data structure:

- **preprocess(n):** Set the universe $U = \{1, \dots, n\}$. Create $k = n$ sets S_1, \dots, S_k where $S_i = \{i\}$. The **preprocess** operation should always be called before any other operation in the data structure; calling it again also will ‘restart’ the data structure (so it should only be called once at the beginning).
- **find(e):** Given an element $e \in U$, return the index of the set S_j where $e \in S_j$ (return j).
- **union(i, j):** Given index of two sets i, j , replace S_i, S_j with $S_i \cup S_j$ (with index of the new set becoming $\min\{i, j\}$).

These three are all the operations we require from our data structure.

Example:

- Suppose we first run **preprocess**(6) which creates $U = \{1, 2, 3, 4, 5, 6\}$ and the following $k = 6$ sets:

$$S_1 = \{1\}, \quad S_2 = \{2\}, \quad S_3 = \{3\}, \quad S_4 = \{4\}, \quad S_5 = \{5\}, \quad S_6 = \{6\}.$$

- Running **union**(2, 3) results in:

$$S_1 = \{1\}, \quad S_2 = \{2, 3\}, \quad S_4 = \{4\}, \quad S_5 = \{5\}, \quad S_6 = \{6\}.$$

- Running **union**(1, 4) results in:

$$S_1 = \{1, 4\}, \quad S_2 = \{2, 3\}, \quad S_5 = \{5\}, \quad S_6 = \{6\}.$$

- At this point:

$$\text{find}(1) = 1, \quad \text{find}(2) = 2, \quad \text{find}(3) = 2, \quad \text{find}(4) = 1, \quad \text{find}(5) = 5, \quad \text{find}(6) = 6.$$

- We can further run **union**(1, 2) which results in:

$$S_1 = \{1, 2, 3, 4\}, \quad S_5 = \{5\}, \quad S_6 = \{6\}.$$

...

Implementation: There are different ways of implementing such a data structure and this problem was studied extensively in 60's, 70's, and 80's and at this point is very well-understood. We will not go over these implementations in this course¹. For our purpose, we only state the following bounds that correspond to the (almost) easiest implementation of this data structure:

- **preprocess**(n) running in $O(n)$ time;
- **find**(e) running in $O(\log n)$ time for every element $e \in U$;
- **union**(i, j) running in $O(\log n)$ time also for every two given indices i, j of the sets.

How to Use Union-Find to Speed-Up Kruskal's Algorithm?

We now implement Kruskal Algorithm using the union-find data structure. The idea is to maintain the connected components of F as sets S_1, \dots, S_k inside the union-find data structure, update them accordingly throughout the algorithm, and use them to find whether an edge creates a cycle or not inside F .

1. Sort the edges of G in increasing (non-decreasing) order of their weights. Let $F = \emptyset$.
2. Create a union-find data structure D with **preprocess**(n) (with universe $U = \{v_1, \dots, v_n\}$).
3. For $i = 1$ to m (in the sorted ordering of edges):
 - (a) Let $e_i = (u_i, v_i)$. Let $a = D.\text{find}(u_i)$ and $b = D.\text{find}(v_i)$.
 - (b) If $a = b$ continue to the next edge.
 - (c) Otherwise, $F \leftarrow F \cup \{e_i\}$ and run $D.\text{union}(a, b)$.
4. Return F .

¹But you are strongly encouraged to check this data structure online (or for instance in Chapter 21 of the CLRS book) for mind boggling aspects of these implementations

To show that this algorithm works we only need to show that we will add e_i to F if and only if F adding e_i does *not* make a cycle – this way, this algorithm will be identical to what described above.

To see how this algorithm works, simply consider the sets S_1, \dots, S_k maintained by the data structure. We prove by induction over index i of the for-loop that S_1, \dots, S_k at any point corresponds to the connected components of F (this is our induction hypothesis). The base case of the induction, say for simplicity when $i = 0$, corresponds to the beginning when $F = \emptyset$; in this case the sets in the data structure D are singleton sets (by definition of **preprocess**) which are exactly the connected components of F .

We now prove the induction step. Suppose up until some iteration i , we have connected components and sets in D equal to S_1, \dots, S_k . We prove that after this iteration also this continues to hold. If we skip the edge e_i , we are neither changing D nor F so connected components of F and sets in D remain the same and thus by induction hypothesis we are correct. If we add e_i to F , we the two components containing endpoints of this edge will become the same component in F (they are now connected); at the same time, running $D.\text{union}$ ensures that the corresponding sets in D will also be merged. Thus, in this case also, the connected components of F correspond to sets in D , proving the induction step.

Finally, note that in the algorithm, whenever we skip an edge e_i , it is because both end points of e_i belong to the same connected component (using above argument): since adding an edge to a connected component *always* makes a cycle, we should indeed skip adding this edge. On the other hand, when we decide to add the edge e_i to F , it is because endpoints of e_i belong to two different connected components (again using above argument): since adding an edge between two different connected components *never* create a cycle, we should indeed add this edge to F . This concludes the proof.

1.2 Prim's Algorithm

We now go over Prim's algorithm for finding an MST. As can be easily observed in Kruskal's algorithm, we are maintaining *multiple* connected component of the forest F simultaneously and update them (merge them together in the algorithm). On the other hand, Prim's algorithm only contains *one non-trivial* connected component (all other components are singleton vertices) which we “grow” until it contains the entire graph. The strategy for growing the component is also very basic: we simply pick the edge with minimum weight that goes out of this component in every step. Prim's algorithm is formally as follows:

Prim's Algorithm:

1. Let $\text{mark}[1 : n] = \text{FALSE}$ and s be an *arbitrary* vertex of the graph.
2. Let $F = \emptyset$.
3. Set $\text{mark}[s] = \text{TRUE}$ and let S (initially) be the set of edges incident on s .
4. While S is non-empty:
 - (a) Let $e = \{u, v\}$ be the *minimum weight* edge in S and remove e from S .
 - (b) If $\text{mark}[u] = \text{mark}[v] = \text{TRUE}$ ignore this edge and go to the next iteration of the while-loop.
 - (c) Otherwise, without loss of generality, assume $\text{mark}[v] = \text{FALSE}$ (if $\text{mark}[u] = \text{FALSE}$ simply switch the name of u, v below).
 - (d) Set $\text{mark}[v] = \text{TRUE}$, add all edges incident on v to S , and add $\{u, v\}$ to F .
5. Return F .

Proof of Correctness: There are two steps in proving the correctness of Prim's algorithm: (1) it outputs a spanning tree, and (2) the weight of this spanning tree is minimum, namely, it is an MST. The proof of first part is identical to the proof of correctness of BFS and DFS (note that the only difference between these algorithms is that in BFS we stored vertices in a queue and removed them while here we are storing the edges and then examine their unmarked endpoints). We thus prove the second (and the main) part below.

The plan as before is to show that Prim's algorithm implements the meta-algorithm discussed earlier by picking safe edge. This in turn is done by proving that any edge $\{u, v\}$ added to F satisfies the properties of Theorem 1 and thus is a safe edge.

At the beginning of each iteration of the while-loop, define C as the set of all marked vertices. By part (1) of the argument (and correctness of DFS/BFS), C is the connected component of s in the forest F at the beginning of this iteration. Consider the cut $(C, V - C)$ in this iteration. This cut has no cut-edges in F by definition as C is a connected component. We claim that the edge e chosen in this iteration, either has both its endpoints marked TRUE, or is the minimum weight edge among the cut-edges of $(C, V - C)$ in G which implies that this edge is a safe edge by Theorem 1. As such, we only add an edge e to F if this edge is safe.

We prove this by showing that *all* cut-edges of $(C, V - C)$ in G belong to the set S at the beginning of this iteration: whenever a vertex v joins C (by setting $mark[v] = TRUE$) we add all edges incident on v to S and we definitely never removed any cut-edge of $(C, V - C)$ before as otherwise we already added those edges to F and C would not be a connected component of F anymore (note that for a cut-edge only one endpoint can be marked TRUE). Since we are returning a minimum weight edge from S at this step, we know that weight of $\{u, v\}$ is smaller than weight of all edges in S and in particular all cut-edges of $(C, V - C)$ in G . This proves that $\{u, v\}$ is a minimum weight edge in the cut $(C, V - C)$.

Runtime Analysis: Similar to Kruskal's algorithm, there is a direct (but not so efficient) way of implementing this algorithm. Simply store the set S in an array or a linked list and then in each iteration of the while-loop, spend $O(|S|) = O(m)$ time to find the minimum weight edge of S . It is easy to see that in this case the total runtime of the algorithm would be $O(n + m + m^2) = O(m^2)$. However, we can implement this algorithm more cleverly in $O(m \log m)$ time using a standard data structure, namely, a *min-heap*.

Detour: Min-Heap Data Structure

We now briefly go over the *min-heap* data structure that you most likely have seen in your previous courses. The goal of this data structure is to maintain a collection of numbers (or more generally *(key, value) pairs*), allow us to add more numbers to this collection, and extract the *minimum* number from it.

We now define the data structure formally. To do so, we simply need to define the operations supported by the data structure:

- **preprocess:** Create an empty set $U = \emptyset$ which will (later on) contain the numbers we insert to the heap. The **preprocess** operation should always be called before any other operation in the data structure; calling it again also will 'restart' the data structure (so it should only be called once).
- **add(e):** Adds the element e to U .
- **extract-min:** Remove the minimum number from U and returns it.
- **size:** returns the number of elements in the heap.

These three are all the operations we require from our data structure (although some variants of min-heap support more functionalities as well).

Implementation: There is a simple way of implementing the min-heap data structure that many of you have probably seen already (see Chapter 6.1 of CLRS for a refresher). For our purpose, we only state the following bounds that follow from this simple implementation:

- **preprocess** runs in $O(1)$ time;
- **add(e)** runs in $O(\log n)$ time where n is the number of the elements in the heap;
- **extract-min** runs in $O(\log n)$ time where n is the number of elements in the heap;
- **size** runs in $O(1)$ time.

How to Use Min-Heap to Speed-Up Prim's Algorithm?

We run the following algorithm by replacing the set S of edges with a min-heap H .

1. Let $mark[1 : n] = FALSE$ and s be an *arbitrary* vertex of the graph.
2. Let $F = \emptyset$ be the maintained forest and set $mark[s] = TRUE$.
3. Let H be a min-heap data structure: Call $H.preprocess$ and $H.add(e)$ for every edge e incident on s .
4. While $H.size > 0$:
 - (a) Let $e = \{u, v\} = H.extract-min$.
 - (b) If $mark[u] = mark[v] = TRUE$ ignore this edge and go to the next iteration of the while-loop.
 - (c) Otherwise, without loss of generality, assume $mark[v] = FALSE$ (if $mark[u] = FALSE$ simply switch the name of u, v below).
 - (d) Set $mark[v] = TRUE$, call $H.add(e)$ for all edges e incident on v , and add $\{u, v\}$ to F .
5. Return F .

It is immediate to verify that this algorithm is the same exact implementation of Prim's algorithm above with the only difference being that the set S replaced by a min-heap H . The runtime of this algorithm is now $O(\log m)$ by each iteration of the while-loop to extract the minimum and $O(\deg(v) \cdot \log m)$ for each vertex v to insert all edges incident to v in the heap (note that size of the min-heap never gets more than $O(m)$). This means that the total runtime is $O(m \log m)$ as desired.