

CS 344: Design and Analysis of Computer Algorithms


(Spring 2022 — Sections 5,6,7,8)

Lecture 22: Applications of Network Flow

The Network Flow Problem


Motivation

- Think of a collection of pipes
- A source **s** and a sink **t** — the source produces some material, say, water, and the sink consumes it
- Each pipe **e** can carry a certain amount of water c_e at any point of time
- Maximize the rate of sending water from source to sink




Networks

- Directed graph $G=(V,E)$
- A source vertex s and a sink vertex t
- Capacity c_e on any edge e



Flow: Example

- A function $f: V \times V \rightarrow \mathbb{R}$ with the following properties:
- **Capacity constraint**
- **Preservation of flow**




Flow

- A function $f: V \times V \rightarrow \mathbb{R}$ with the following properties:

- **Capacity constraint**

- **Preservation of flow**

- **Value** of a flow f : amount of flow leaving $s = \sum_{v \in V} f(s, v)$



Network Flow Problem

- Network flow (or maximum flow) problem:
- **Input:**
 - A network $G=(V,E)$ with edge-capacities and a source and a sink
- **Output:**
 - Find a flow with **largest value** in G

Ford-Fulkerson


- For concreteness, we stick with the first and simplest network flow algorithm
- **Ford-Fulkerson Algorithm:**
 - Solves maximum flow in $O(m \cdot F)$ time where F is the value of maximum flow from s to t .

Application I: Edge- Disjoint Paths

Application II: Bipartite Matching


Bipartite Matching

- Bipartite graph $G=(V,E)$:
 - There are two sets of vertices L and R
 - All edges are only between L and R




Bipartite Matching

- Bipartite graph $G=(V,E)$:
 - There are two sets of vertices L and R
 - All edges are only between L and R
- Matching M in G :
 - A subset of edges in E
 - No vertex used more than once




Bipartite Matching

- **Bipartite** graph $G=(V,E)$:
 - There are two sets of vertices **L** and **R**
 - All edges are only **between** **L** and **R**
- **Matching** **M** in **G**:
 - A subset of edges in **E**
 - No vertex used **more than once**




Bipartite Matching

- **Bipartite** graph $G=(V,E)$:
 - There are two sets of vertices **L** and **R**
 - All edges are only **between** **L** and **R**
- **Matching** **M** in **G**:
 - A subset of edges in **E**
 - No vertex used **more than once**




Bipartite Matching

- Bipartite graph $G=(V,E)$:
 - There are two sets of vertices L and R
 - All edges are only between L and R
- Matching M in G :
 - A subset of edges in E
 - No vertex used more than once



Bipartite Matching

- Bipartite graph $G=(V,E)$:
 - There are two sets of vertices L and R
 - All edges are only between L and R
- Matching M in G :
 - A subset of edges in E
 - No vertex used more than once




Bipartite Matching Problem

- **Input:**
 - a bipartite graph $G=(V,E)$ with bipartition L and R
- **Output:**
 - Output a maximum matching in G , i.e., a matching with the largest number of edges


Bipartite Matching Problem

- **Input:**
 - a **bipartite** graph $G=(V,E)$ with bipartition **L** and **R**
- **Output:**
 - Output a **maximum matching** in **G**, i.e., a matching with the **largest number of edges**



Bipartite Matching Problem


- **Input:**
 - a **bipartite** graph $G=(V,E)$ with bipartition **L** and **R**
- **Output:**
 - Output a **maximum matching** in **G**, i.e., a matching with the **largest number of edges**





Bipartite Matching Problem

- Applications:
 - Online advertising
 - Auctions and markets
 - Students-Dorms Assignments
 - Kidney exchange program
 - ...



Reduction to Network Flow





Reduction to Network Flow





Reduction to Network Flow



Reduction to Network Flow



Reduction to Network Flow



Reduction to Network Flow

- Create a network $G'=(V',E')$ as follows:
 - Copy the vertices in L and R of G in G' also
 - For any edge (u,v) in G with u in L and v in R , add a directed edge from u to v in G' . Set the capacities to $+\infty$
 - Add new vertices s and t , which will be source and sink
 - Connect s to every vertex in L with capacity 1
 - Connect every vertex in R to t with capacity 1
- Compute a maximum flow f in G'
- Return edges (u,v) in G if u in L and v in R and $f(u,v) = 1$

Proof of Correctness

- Create a network $G'=(V',E')$ as follows:
 - Copy the vertices in L and R of G in G'
 - For any edge (u,v) in G with u in L and v in R , add a **directed edge** from u to v in G' . Set the capacities to $+\infty$
 - Add a **source** s and **sink** t
 - Connect s to vertices in L with capacity 1
 - Connect vertices in R to t with capacity 1
- Compute a maximum flow f in G'
- Return edges (u,v) in G if u in L and v in R and $f(u,v) = 1$
- Part one: A **flow** f of value k gives a **matching** M of size k

Proof of Correctness

- Create a network $G'=(V',E')$ as follows:
 - Copy the vertices in L and R of G in G'
 - For any edge (u,v) in G with u in L and v in R , add a **directed edge** from u to v in G' . Set the capacities to $+\infty$
 - Add a **source** s and **sink** t
 - Connect s to vertices in L with capacity 1
 - Connect vertices in R to t with capacity 1
- Compute a maximum flow f in G'
- Return edges (u,v) in G if u in L and v in R and $f(u,v) = 1$
- Part two: A **matching** M of size k gives a **flow** f of value k

Proof of Correctness

- Create a network $G'=(V',E')$ as follows:
 - Copy the vertices in L and R of G in G'
 - For any edge (u,v) in G with u in L and v in R , add a directed edge from u to v in G' . Set the capacities to $+\infty$
 - Add a source s and sink t
 - Connect s to vertices in L with capacity 1
 - Connect vertices in R to t with capacity 1
- Compute a maximum flow f in G'
- Return edges (u,v) in G if u in L and v in R and $f(u,v) = 1$
- So the maximum flow f gives a maximum matching M

Runtime Analysis

- Create a network $G'=(V',E')$ as follows:
 - Copy the vertices in L and R of G in G'
 - For any edge (u,v) in G with u in L and v in R , add a directed edge from u to v in G' . Set the capacities to $+\infty$
 - Add a source s and sink t
 - Connect s to vertices in L with capacity 1
 - Connect vertices in R to t with capacity 1
- Compute a maximum flow f in G'
- Return edges (u,v) in G if u in L and v in R and $f(u,v) = 1$
- Creating G' takes $O(n+m)$ time
- G' has $n+2$ vertices and $m+2n$ edges
- So Ford-Fulkerson takes $O(m \cdot F)$ time

Runtime Analysis

- Create a network $G'=(V',E')$ as follows:
 - Copy the vertices in L and R of G in G'
 - For any edge (u,v) in G with u in L and v in R , add a directed edge from u to v in G' . Set the capacities to $+\infty$
 - Add a source s and sink t
 - Connect s to vertices in L with capacity 1
 - Connect vertices in R to t with capacity 1
- Compute a maximum flow f in G'
- Return edges (u,v) in G if u in L and v in R and $f(u,v) = 1$
- Creating G' takes $O(n+m)$ time
- G' has $n+2$ vertices and $m+2n$ edges
- So Ford-Fulkerson takes $O(m \cdot F)$ time
- $F \leq n/2$ as F is equal to the maximum matching size and that is $\leq n/2$
- So it takes $O(mn)$ time

Application III: Exam Scheduling

Exam Scheduling Problem

- **Input:**

- c courses with course i having $E[i]$ enrolled students
- r rooms with room j having $S[j]$ available seats
- t available time-slots for exams denoted by $\{1, 2, \dots, t\}$
- p proctors with proctor k being available at times $T[k] \subseteq \{1, \dots, t\}$

Exam Scheduling Problem

- **Input:**

- c courses with course i having $E[i]$ enrolled students
- r rooms with room j having $S[j]$ available seats
- t available time-slots for exams denoted by $\{1, 2, \dots, t\}$
- p proctors with proctor k being available at times $T[k] \subseteq \{1, \dots, t\}$

- **Constraints:**

- Room j can only be assigned to course i if $E[i] \leq S[j]$
- Any room can be assigned to only one exam in a given time-slot
- A proctor k can only attend exams at time-slots in $T[k]$
- No proctor can work for more than 3 exams and each exam needs exactly 1 proctor

Exam Scheduling Problem

- **Input:**

- c courses with course i having $E[i]$ enrolled students
- r rooms with room j having $S[j]$ available seats
- t available time-slots for exams denoted by $\{1, 2, \dots, t\}$
- p proctors with proctor k being available at times $T[k] \subseteq \{1, \dots, t\}$


- **Constraints:**

- Room j can only be assigned to course i if $E[i] \leq S[j]$
- Any room can be assigned to only one exam in a given time-slot
- A proctor k can only attend exams at time-slots in $T[k]$
- No proctor can work for more than 3 exams and each exam needs exactly 1 proctor


- **Output:**

- Output c tuples (course, room, time-slot, proctor) satisfying the constraints or say it is not possible

Reduction to Network Flow



Reduction to Network Flow



Reduction to Network Flow

- Create $G = (V, E)$ with 4 layers of vertices C , R , T , and P and one source s and one sink t
 - c vertices corresponding to courses in C
 - r vertices corresponding to rooms in R
 - t vertices corresponding to time-slots in T
 - p vertices corresponding to proctors in P


Reduction to Network Flow

- Create $G = (V, E)$ with 4 layers of vertices C , R , T , and P and one source s and one sink t
- Add the following edges:
 - Source s to all vertices in C with capacity 1
 - Vertex i in C to vertex j in R if $E[i] \leq S[j]$ (we can hold exam of course i in room j) with capacity 1
 - All vertices in R to all vertices in T with capacity 1
 - Vertex t in T to vertex k in P if t in $T[k]$ with capacity 1 (proctor k can work on an exam at time t)
 - Vertex k in P to t with capacity 3


Reduction to Network Flow

- Create $G = (V, E)$ with 4 layers of vertices C , R , T , and P and one source s and one sink t
- Add the following edges
- Find a maximum flow f in the network G
- For any flow-path $(s, i \text{ in } C, j \text{ in } R, t \text{ in } T, p \text{ in } P, t)$, we return the tuple (i, j, t, p)
- If the number of tuples is less than c , we say scheduling is not possible

Reduction to Network Flow



Reduction to Network Flow



Proof of Correctness

- Part one: any flow f of value ℓ gives a valid schedule of ℓ courses:
- So we output ℓ valid tuples in the schedule

Proof of Correctness

- Part two: any valid schedule of ℓ courses gives a flow of value ℓ
- We thus have maximum flow f will give largest valid schedule

Runtime Analysis

- Network G has $n = c + r + t + p + 2$ vertices
- So it has at most $m = O(n^2)$ edges
- It takes $O(n + m) = O(n^2)$ times to create the network
- Moreover, largest possible flow has value $F \leq c \leq n$
- So Ford-Fulkerson takes $O(mF) = O(n^3)$ time